

CLab-4 Report

Jeff Yuanbo Han u6617017

May 13, 2018

In computer lab 4, I experiment both tasks. And My major task is CNN (deep learning) for MNIST handwritten digit recognition; Task 2—Camera calibration will serve as a bonus.

Contents

1	CNN Based Vision Recognition	1
1.1	My CNN Model	1
1.2	Batch Size	2
1.3	Learning Rate	3
1.4	Initializer	3
1.5	Optimizer	4
1.6	Drop Out	4
1.7	Hidden Layer	5
2	Camera Calibration	6
2.1	Illustration by Pictures	6
2.2	Calibration Data	6
A	Python & MATLAB Codes	8
A.1	CNN Based Vision Recognition	8
A.1.1	<i>cnn_1.1.py</i>	8
A.1.2	<i>cnn_1.2.py</i>	10
A.1.3	<i>cnn_2.1.py</i>	12
A.2	Camera Calibration	15
A.2.1	<i>calibrate.m</i>	15
A.2.2	<i>selectPoints.m</i>	16
A.2.3	<i>estimateC.m</i>	17

1 CNN Based Vision Recognition

1.1 My CNN Model

To achieve the goal of recognizing MNIST handwritten digits, I have tried tens of neural networks, ranging from the basic single-layer net to the classic LeNet-5. However, the most effective architectures usually involves more than 3 hidden layers, and therefore is quite time-consuming. The guideline of CLab-4 requires our model to limit training procedure to 3 minutes. So I have to prune those subtle but complex structures. After myriads of experiments, I finally choose my deep neural network as the one contain a convolutional layer and a full-connected layer:




Figure 1: Benchmark: my final model

Layer	Window Size	Features/Units
Input	28×28	1
Convolution	5×5	6
Max-pool	2×2	6
Full-connected	1	300
Output	1	10

All the activation functions are using ReLU, i.e. Rectified Linear Unit, mathematically, $\max(0, input)$. As a result of experiments, ReLU is the most efficient among functions that are commonly used for activation (Softmax, Tanh, etc.). This means performing, say Softmax, in a 2-hidden-layer network can hardly complete training under 3 minutes.

My Settings are

$$\begin{aligned}
&\text{Initializer} = \text{Trimmed Normal Distribution } (\mu = 0, \sigma = 0.3); \\
&\text{Batch size} = 50; \\
&\text{Iteration} = 5000; \\
&\text{Learning rate} = 0.001; \\
&\text{Optimizer} = \text{Adam};
\end{aligned} \tag{1}$$

The training time is on average 175s. And the test accuracy is usually above 90%. With several attempts, the best run arrives at 95.33%. This result can still be better with larger batch-size, but it will violate the rule of time. The training error descent procedure is shown below in Figure 1. In following sections, you will see each hyper-parameter of my settings is at least a local optima to some extent, and they overall account for my choice of model.

1.2 Batch Size

In this section, I experiment with the batch-size. Maintaining my model architecture and all the other hyper-parameters except batch-size and iteration. When batch-size gets bigger, the epochs, which reflects the real amount of computation, will increase meanwhile. So we may have to reduce the number of iterations in order to satisfy the 3-min requirement, and vice versa (can train more steps with small batch-size). The Batch Size Table and Figure-2 show an approximate comparison of different batch-sizes.




Figure 2: Batch size = 10 & 200




Figure 3: Learning rate = 0.1 & 0.0001

Batch size	Iteration	Test accuracy
10	6500	52%
50	5000	95%
100	3500	88%
200	2000	79%

1.3 Learning Rate

Experimenting with learning rate is much easier. Just keep the Baseline settings and change learning rate only. Results are displayed in the Learning Rate Table and Figure . As a consequence, we could infer that learning rate ≥ 0.01 is too big for parameters to converge, and ≤ 0.001 is so small that the model has not been trained adequately.

Learning rate	Test accuracy
0.1	15%
0.01	47%
0.001	95%
0.0001	41%

1.4 Initializer

Trying to change other initializers also get different outputs. Please look at the Initializer Table and Figure 4.




Figure 4: Initializer = Uniform(-1, 1) & Constant 0




Figure 5: Optimizer = Adagrad & SGD

Distribution	Parameter	Test accuracy
Trimmed Normal	$\mu = 0, \sigma = 0.1$	82%
Trimmed Normal	$\mu = 0, \sigma = 0.3$	95%
Trimmed Normal	$\mu = 0, \sigma = 0.5$	84%
Uniform	$[-1, 1]$	18%
Point (Constant)	0	88%

1.5 Optimizer

Similarly as before, remain all the other setting as Baseline, and then try different optimizers (see the Optimizer Table and Figure 5). We find that other optimizers do not agree with this specific model, possibly in that all the hyper-parameters are set and fixed under Adam optimizer.

Optimizer	Test accuracy
Adam	95%
Adagrad	24%
SGD	17%

1.6 Drop Out

I am not using drop-out technique in my final model, but I do have studied on it. With drop-out rate = 0.5, the model can achieve 76% or so test accuracy (see the Optimizer Table and Figure 6).




Figure 6: Drop-out rate = 0 & 0.5




Figure 7: *cnn_1.2* & *cnn_2.1*

In other words, the result has not been improved at least within 3 minutes; instead, it still needs a number of training processes. This is probably because our network has such an elementary topology that adopting dropout is needless.

Drop-out rate	Test accuracy
0	95%
0.5	76%

1.7 Hidden Layer

As described in Section 1.1, my final model is *cnn_1.1*, i.e. 1 convolutional layer (and max-pooling) + 1 full-connected layer. I have also tried 2 convolutional layers (*cnn_2.1*) or 2 full-connected layers (*cnn_1.2*). They of course have the potential to exceed *cnn_1.1* theoretically. However, this is a time-limited task, so we have to cut the training iterations to accord with 3-min requirement. Under this circumstance, results are show below in the Model Table and Figure 7.

Model	Convolution + Max-pool	Full-connected	Iteration	Test accuracy
<i>cnn_1.1</i>	6 features	300	5000	95%
<i>cnn_1.2</i>	6 features	300 + 180	3500	80%
<i>cnn_2.1</i>	6 + 16 features	180	3000	35%

In addition, *cnn_2.2* becomes the classic LeNet-5[1], which has obtained 99.05% test accuracy, and can be better with appropriate distorted training images. However, such complicated networks are impossible to run out within 3 minutes.

2 Camera Calibration

2.1 Illustration by Pictures

My implement of estimating Calibration matrix is listed at *calibrate.m*. I have selected two images to check the Calibration, namely: stereo2012a.jpg and stereo2012d.jpg (Figure 8). In these 2 pictures, red crosses are the points I choose originally, whereas blue circles are the estimated locations of the chosen points projected by Calibration. The green lines show the projected X, Y, Z axes.

2.2 Calibration Data

Take stereo2012d.jpg as an example. The outputs after running *selectPoints.m* and *estimateC.m* are:

MSE = 0.134857

C =

-0.0052	0.0022	0.0145	-0.7640
-0.0013	0.0141	-0.0006	-0.6448
0.0000	0.0000	0.0000	-0.0020

K =

805.6066	8.3036	414.9998
0	795.2093	216.8355
0	0	1.0000

R =

0.7064	0.0312	-0.7071
0.2711	-0.9348	0.2296
-0.6538	-0.3539	-0.6688

t =

76.1529
55.9128
71.1172

Let f_x = horizontal focal length, f_y = vertical focal length, then

$$\begin{aligned}f_x &= K[1, 1] = 805.6066 \\f_y / \sin \theta &= K[2, 2] = 795.2093 \\-f_x \cot \theta &= K[1, 2] = 8.3036\end{aligned}$$

Solving them, we derive

$$\begin{aligned}f_x &= 805.6066 \\f_y &= 795.1671\end{aligned}$$





Figure 8: stereo2012a.jpg & stereo2012d.jpg

Besides, it is not hard to know

$$\sin \alpha_y = R[1, 3] = -0.7071$$

Thus we obtain

$$\alpha_y = -0.7854 \text{ rad} \approx -45^\circ \text{ degree}$$

which is the pitch angle of camera with respect to X-Z plane in the world coordinate system.

References

- [1] Lecun, Y., Bottou, L., Bengio, Y., & Haffner, P. (1998). Gradient-based learning applied to document recognition. Proceedings of the IEEE, 86(11), 2278-2324.

A Python & MATLAB Codes

A.1 CNN Based Vision Recognition

A.1.1 *cnn_1_1.py*

```
1  # -*- coding: utf-8 -*-
2  """
3  1 * (Convolution + Max_pool) + 1 * Full-connected
4  """
5  import tensorflow as tf
6  from tensorflow.examples.tutorials.mnist import input_data
7  from random import sample
8  import matplotlib.pyplot as plt
9  from time import time
10
11
12  def weight_variable(shape):
13      initial = tf.truncated_normal(shape, stddev=0.3)
14      #initial = tf.random_uniform(shape=shape, minval=-1, maxval=1)
15      return tf.Variable(initial)
16
17
18  def bias_variable(shape):
19      #initial = tf.constant(0.1, shape=shape)
20      #initial = tf.random_uniform(shape=shape, minval=-1, maxval=1)
21      initial = tf.truncated_normal(shape, stddev=0.3)
22      return tf.Variable(initial)
23
24
25  def conv2d(x, W):
26      return tf.nn.conv2d(x, W, strides=[1, 1, 1, 1], padding='VALID')
27
28
29  def max_pool_2x2(x):
30      return tf.nn.max_pool(x, ksize=[1, 2, 2, 1],
31      strides=[1, 2, 2, 1], padding='VALID')
32
```



```

33
34 if __name__ == '__main__':
35     # Read in MNIST data
36     mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
37     ind = sample(range(mnist.train.images.shape[0]), 10000)
38     train = tf.contrib.learn.datasets.mnist.DataSet(
39         mnist.train.images[ind,:], mnist.train.labels[ind,:], one_hot=True, reshape=False)
40
41     # Placeholders of training attributes and labels
42     x = tf.placeholder(tf.float32, [None, 784])
43     y_ = tf.placeholder(tf.float32, [None, 10])
44
45     # Reshape images from 728*1 to 28*28
46     x_image = tf.reshape(x, [-1, 28, 28, 1])
47
48     # Convolutional layer
49     W_conv1 = weight_variable([5, 5, 1, 6])
50     b_conv1 = bias_variable([6])
51     h_conv1 = tf.nn.relu(conv2d(x_image, W_conv1) + b_conv1)
52     h_pool1 = max_pool_2x2(h_conv1)
53
54     # Full-connected layer, ReLU
55     W_fc1 = weight_variable([12 * 12 * 6, 300])
56     b_fc1 = bias_variable([300])
57     h_pool1_flat = tf.reshape(h_pool1, [-1, 12 * 12 * 6])
58     h_fc1 = tf.nn.relu(tf.matmul(h_pool1_flat, W_fc1) + b_fc1)
59
60     # Placeholder of keep_prob for dropout
61     keep_prob = tf.placeholder(tf.float32)
62
63     # Output layer with 10 classes
64     W_fc2 = weight_variable([300, 10])
65     b_fc2 = bias_variable([10])
66     h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)
67     y_conv = tf.matmul(h_fc1_drop, W_fc2) + b_fc2
68
69     # Compute cost: Cross entropy after softmax
70     cross_entropy = tf.reduce_mean(
71         tf.nn.softmax_cross_entropy_with_logits(labels=y_, logits=y_conv))
72
73     # Learning step, optimized by Adam
74     train_step = tf.train.AdamOptimizer(1e-3, 0).minimize(cross_entropy)
75
76     # Define accuracy
77     correct_prediction = tf.equal(tf.argmax(y_conv, 1), tf.argmax(y_, 1))
78     accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
79
80     # Create Session and initialize variables
81     sess = tf.InteractiveSession()
82     sess.run(tf.global_variables_initializer())
83

```

```

84 # Training process
85 iteration = 5000
86 err = []
87 start_time = time()
88 for i in range(iteration):
89     batch = train.next_batch(50)
90     # Report the accuracy on training set every 100 steps
91     if i % 100 == 0:
92         train_accuracy = accuracy.eval(feed_dict={
93             x: train.images, y_: train.labels, keep_prob: 1.0})
94         print("step %d, train accuracy %g" % (i, train_accuracy))
95         err.append(1-train_accuracy)
96         train_step.run(feed_dict={x: batch[0], y_: batch[1], keep_prob: 1.0})
97
98     end_time = time()
99     print('Training time: {}s'.format(end_time-start_time))
100
101 # Accuracy on test set
102 print("test accuracy %g" % accuracy.eval(feed_dict={
103     x: mnist.test.images, y_: mnist.test.labels, keep_prob: 1.0}))
104
105 # Plot training error
106 plt.plot(range(len(err)), err)
107 plt.title('Train Error')
108 plt.xlabel(r'$\times 100$ Iterations')
109 plt.ylabel('Error Rate (%)')
110 plt.xticks([i*iteration/1000 for i in range(11)])
111 plt.yticks([e*0.05 for e in range(21)], [e*5 for e in range(21)])
112 plt.show()

```

A.1.2 cnn_1_2.py

```

1 # -*- coding: utf-8 -*-
2 """
3 1 * (Convolution + Max_pool) + 2 * Full-connected
4 """
5 import tensorflow as tf
6 from tensorflow.examples.tutorials.mnist import input_data
7 from random import sample
8 import matplotlib.pyplot as plt
9 from time import time
10
11
12 def weight_variable(shape):
13     initial = tf.truncated_normal(shape, stddev=0.3)
14     #initial = tf.random_uniform(shape=shape, minval=-1, maxval=1)
15     return tf.Variable(initial)
16
17
18 def bias_variable(shape):
19     #initial = tf.constant(0.1, shape=shape)
20     #initial = tf.random_uniform(shape=shape, minval=-1, maxval=1)

```

```

21 initial = tf.truncated_normal(shape, stddev=0.3)
22 return tf.Variable(initial)
23
24
25 def conv2d(x, W):
26 return tf.nn.conv2d(x, W, strides=[1, 1, 1, 1], padding='VALID')
27
28
29 def max_pool_2x2(x):
30 return tf.nn.max_pool(x, ksize=[1, 2, 2, 1],
31 strides=[1, 2, 2, 1], padding='VALID')
32
33
34 if __name__ == '__main__':
35 # Read in MNIST data
36 mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
37 ind = sample(range(mnist.train.images.shape[0]), 10000)
38 train = tf.contrib.learn.datasets.mnist.DataSet(
39 mnist.train.images[ind,:], mnist.train.labels[ind,:], one_hot=True, reshape=False)
40
41 # Placeholders of training attributes and labels
42 x = tf.placeholder(tf.float32, [None, 784])
43 y_ = tf.placeholder(tf.float32, [None, 10])
44
45 # Reshape images from 728*1 to 28*28
46 x_image = tf.reshape(x, [-1, 28, 28, 1])
47
48 # Convolutional layer
49 W_conv1 = weight_variable([5, 5, 1, 6])
50 b_conv1 = bias_variable([6])
51 h_conv1 = tf.nn.relu(conv2d(x_image, W_conv1) + b_conv1)
52 h_pool1 = max_pool_2x2(h_conv1)
53
54 # Full-connected layer 1, ReLU
55 W_fc1 = weight_variable([12 * 12 * 6, 300])
56 b_fc1 = bias_variable([300])
57 h_pool1_flat = tf.reshape(h_pool1, [-1, 12 * 12 * 6])
58 h_fc1 = tf.nn.relu(tf.matmul(h_pool1_flat, W_fc1) + b_fc1)
59
60 # Placeholder of keep_prob for dropout
61 keep_prob = tf.placeholder(tf.float32)
62
63 # Full-connected layer 2, ReLU
64 W_fc2 = weight_variable([300, 180])
65 b_fc2 = bias_variable([180])
66 h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)
67 h_fc2 = tf.nn.relu(tf.matmul(h_fc1_drop, W_fc2) + b_fc2)
68
69 # Output layer with 10 classes
70 W_fc3 = weight_variable([180, 10])
71 b_fc3 = bias_variable([10])

```

```

72 h_fc2_drop = tf.nn.dropout(h_fc2, keep_prob)
73 y_conv = tf.matmul(h_fc2_drop, W_fc3) + b_fc3
74
75 # Compute cost: Cross entropy after softmax
76 cross_entropy = tf.reduce_mean(
77     tf.nn.softmax_cross_entropy_with_logits(labels=y_, logits=y_conv))
78
79 # Learning step, optimized by Adam
80 train_step = tf.train.AdamOptimizer(1e-3, 0).minimize(cross_entropy)
81
82 # Define accuracy
83 correct_prediction = tf.equal(tf.argmax(y_conv, 1), tf.argmax(y_, 1))
84 accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
85
86 # Create Session and initialize variables
87 sess = tf.InteractiveSession()
88 sess.run(tf.global_variables_initializer())
89
90 # Training process
91 iteration = 3500
92 err = []
93 start_time = time()
94 for i in range(iteration):
95     batch = train.next_batch(50)
96     # Report the accuracy on training set every 100 steps
97     if i % 100 == 0:
98         train_accuracy = accuracy.eval(feed_dict={
99             x: train.images, y_: train.labels, keep_prob: 1.0})
100         print("step %d, train accuracy %g" % (i, train_accuracy))
101         err.append(1-train_accuracy)
102         train_step.run(feed_dict={x: batch[0], y_: batch[1], keep_prob: 1.0})
103
104     end_time = time()
105     print('Training time: {}s'.format(end_time-start_time))
106
107 # Accuracy on test set
108     print("test accuracy %g" % accuracy.eval(feed_dict={
109         x: mnist.test.images, y_: mnist.test.labels, keep_prob: 1.0}))
110
111 # Plot training error
112     plt.plot(range(len(err)), err)
113     plt.title('Train Error')
114     plt.xlabel(r'$\times 100$ Iterations')
115     plt.ylabel('Error Rate (%)')
116     plt.xticks([i*iteration/1000 for i in range(11)])
117     plt.yticks([e*0.05 for e in range(21)], [e*5 for e in range(21)])
118     plt.show()

```

A.1.3 cnn_2_1.py

```

1 # -*- coding: utf-8 -*-
2 """

```

```

3  2 * (Convolution + Max_pool) + 1 * Full-connected
4  """
5  import tensorflow as tf
6  from tensorflow.examples.tutorials.mnist import input_data
7  from random import sample
8  import matplotlib.pyplot as plt
9  from time import time
10
11
12  def weight_variable(shape):
13      initial = tf.truncated_normal(shape, stddev=0.3)
14      #initial = tf.random_uniform(shape=shape, minval=-1, maxval=1)
15      return tf.Variable(initial)
16
17
18  def bias_variable(shape):
19      #initial = tf.constant(0.1, shape=shape)
20      #initial = tf.random_uniform(shape=shape, minval=-1, maxval=1)
21      initial = tf.truncated_normal(shape, stddev=0.3)
22      return tf.Variable(initial)
23
24
25  def conv2d(x, W):
26      return tf.nn.conv2d(x, W, strides=[1, 1, 1, 1], padding='VALID')
27
28
29  def max_pool_2x2(x):
30      return tf.nn.max_pool(x, ksize=[1, 2, 2, 1],
31          strides=[1, 2, 2, 1], padding='VALID')
32
33
34  if __name__ == '__main__':
35      # Read in MNIST data
36      mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
37      ind = sample(range(mnist.train.images.shape[0]), 10000)
38      train = tf.contrib.learn.datasets.mnist.DataSet(
39          mnist.train.images[ind,:], mnist.train.labels[ind,:], one_hot=True, reshape=False)
40
41      # Placeholders of training attributes and labels
42      x = tf.placeholder(tf.float32, [None, 784])
43      y_ = tf.placeholder(tf.float32, [None, 10])
44
45      # Reshape images from 728*1 to 28*28
46      x_image = tf.reshape(x, [-1, 28, 28, 1])
47
48      # Convolutional layer 1
49      W_conv1 = weight_variable([5, 5, 1, 6])
50      b_conv1 = bias_variable([6])
51      h_conv1 = tf.nn.relu(conv2d(x_image, W_conv1) + b_conv1)
52      h_pool1 = max_pool_2x2(h_conv1)
53

```

```

54 # Convolutional layer 2
55 W_conv2 = weight_variable([5, 5, 6, 16])
56 b_conv2 = bias_variable([16])
57 h_conv2 = tf.nn.relu(conv2d(h_pool1, W_conv2) + b_conv2)
58 h_pool2 = max_pool_2x2(h_conv2)
59
60 # Full-connected layer, ReLU
61 W_fc1 = weight_variable([4 * 4 * 16, 180])
62 b_fc1 = bias_variable([180])
63 h_pool2_flat = tf.reshape(h_pool2, [-1, 4 * 4 * 16])
64 h_fc1 = tf.nn.relu(tf.matmul(h_pool2_flat, W_fc1) + b_fc1)
65
66 # Placeholder of keep_prob for dropout
67 keep_prob = tf.placeholder(tf.float32)
68
69 # Output layer with 10 classes
70 W_fc2 = weight_variable([180, 10])
71 b_fc2 = bias_variable([10])
72 h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)
73 y_conv = tf.matmul(h_fc1_drop, W_fc2) + b_fc2
74
75 # Compute cost: Cross entropy after softmax
76 cross_entropy = tf.reduce_mean(
77     tf.nn.softmax_cross_entropy_with_logits(labels=y_, logits=y_conv))
78
79 # Learning step, optimized by Adam
80 train_step = tf.train.AdamOptimizer(1e-3, 0).minimize(cross_entropy)
81
82 # Define accuracy
83 correct_prediction = tf.equal(tf.argmax(y_conv, 1), tf.argmax(y_, 1))
84 accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
85
86 # Create Session and initialize variables
87 sess = tf.InteractiveSession()
88 sess.run(tf.global_variables_initializer())
89
90 # Training process
91 iteration = 3000
92 err = []
93 start_time = time()
94 for i in range(iteration):
95     batch = train.next_batch(50)
96     # Report the accuracy on training set every 100 steps
97     if i % 100 == 0:
98         train_accuracy = accuracy.eval(feed_dict={
99             x: train.images, y_: train.labels, keep_prob: 1.0})
100         print("step %d, train accuracy %g" % (i, train_accuracy))
101         err.append(1-train_accuracy)
102         train_step.run(feed_dict={x: batch[0], y_: batch[1], keep_prob: 1.0})
103
104     end_time = time()

```

```

105 print( 'Training time: {}s '.format(end_time-start_time))
106
107 # Accuracy on test set
108 print("test accuracy %g" % accuracy.eval(feed_dict={
109 x: mnist.test.images, y_: mnist.test.labels, keep_prob: 1.0}))
110
111 # Plot training error
112 plt.plot(range(len(err)), err)
113 plt.title( 'Train Error ')
114 plt.xlabel(r '$\times 100$ Iterations ')
115 plt.ylabel( 'Error Rate (%) ')
116 plt.xticks([i*iteration/1000 for i in range(11)])
117 plt.yticks([e*0.05 for e in range(21)], [e*5 for e in range(21)])
118 plt.show()

```

A.2 Camera Calibration

A.2.1 *calibrate.m*

```

1 function C = calibrate(im, XYZ, uv)
2 % Function to perform camera calibration.
3 %
4 % Usage:    K = calibrate(image, XYZ, uv)
5 %
6 %   Where:   image - is the image of the calibration target.
7 %             XYZ - is a N x 3 array of XYZ coordinates
8 %                   of the calibration target points.
9 %             uv  - is a N x 2 array of the image coordinates
10 %                  of the calibration target points.
11 %             K   - is the 3 x 4 camera calibration matrix.
12 % The variable N should be an integer greater than or equal to 6.
13 %
14 % This function plots the uv coordinates onto the image of the calibration
15 % target.
16 %
17 % It also projects the XYZ coordinates back into image coordinates using
18 % the calibration matrix and plots these points too as
19 % a visual check on the accuracy of the calibration process.
20 %
21 % Lines from the origin to the vanishing points in the X, Y and Z
22 % directions are overlaid on the image.
23 %
24 % The mean squared error between the positions of the uv coordinates
25 % and the projected XYZ coordinates is also reported.
26 %
27 % The function should also report the error in satisfying the
28 % camera calibration matrix constraints.
29 %
30 % By Jeff Yuanbo Han (u6617017), 2018-05-13.
31
32 n = size(uv, 1); % number of points
33 A = zeros(2*n, 12);

```

```

34 for i = 1:n
35 A(2*i-1, 5:8) = -[XYZ(i,:), 1];
36 A(2*i-1, 9:12) = uv(i,2) * [XYZ(i,:), 1];
37 A(2*i, 1:4) = [XYZ(i,:), 1];
38 A(2*i, 9:12) = -uv(i,1) * [XYZ(i,:), 1];
39 end
40
41 V] = svd(A);
42 C = reshape(V(:,end), [4,3])';
43
44 err = 0; % Squared error
45
46 % Display the selected and estimated points
47 figure; imshow(im); hold on;
48 for i = 1:n
49 plot(uv(i,1), uv(i,2), 'rx');
50 estim = C * [XYZ(i,:),1]';
51 estim = estim ./ estim(3);
52 plot(estim(1), estim(2), 'bo');
53
54 err = err + dist(uv(i,:), estim
55 end
56
57 fprintf('MSE = %f\n', err/n);
58
59 orig = C * [0,0,0,1]';
60 orig = orig ./ orig(3);
61 x_axis = C * [50,0,0,1]';
62 x_axis = x_axis ./ x_axis(3);
63 y_axis = C * [0,50,0,1]';
64 y_axis = y_axis ./ y_axis(3);
65 z_axis = C * [0,0,50,1]';
66 z_axis = z_axis ./ z_axis(3);
67 plot([orig(1),x_axis(1)], [orig(2),x_axis(2)], 'g');
68 plot([orig(1),y_axis(1)], [orig(2),y_axis(2)], 'g');
69 plot([orig(1),z_axis(1)], [orig(2),z_axis(2)], 'g');
70 end

```

A.2.2 selectPoints.m

```

1 % CLAB-4: Select points for estimating.
2 % By Jeff Yuanbo Han (u6617017), 2018-05-13.
3 img = imread('stereo2012d.jpg');
4
5 imshow(img);
6 display('click mouse for 6 features...')
7 uv = ginput(6); % Graphical user interface to get 12 points
8 display(uv);
9
10 XYZ = [ 7, 7, 0;
11 0, 7, 7;
12 7, 0, 7;

```



```

13 14, 7, 0;
14 0,14, 7;
15 7, 0,14 ];
16
17 save d_points.mat img uv XYZ

```

A.2.3 *estimateC.m*

```

1 % CLAB-4: Estimate and decompose C.
2 % By Jeff Yuanbo Han (u6617017), 2018-05-13.
3 load d_points.mat
4 C = calibrate(img, XYZ, uv)
5 [K, R, t] = vgg_KR_from_P(C)
6
7 save d_points.mat C K R t -append

```